

MLJ Workshop

Resources: github.com/ablaom/MLJTutorial

Anthony Blaom with Quinn Asena

The Alan Turing Institute



THE UNIVERSITY OF
AUCKLAND
Te Whare Wānanga o Tāmaki Makaurau
NEW ZEALAND



The Team

Core design: A. B., Franz Kiraly, Sebastian Vollmer

Lead development: A. B., Thibaut Lienart

Other contributors: Diego Arenas, Samuel Okon, Yiannis Simillides, Julian Samaroo, Ayush Shridar, Geoffroy Dolphin, Mosè Giordano...

Julia language consultants: Avik Sengupta

Some machine learning libraries in Julia

```
22-element Vector{String}:
"OutlierDetectionNeighbors"
"OutlierDetectionPython"
"OutlierDetectionNetworks"
"ScikitLearn"
"DecisionTree"
"MultivariateStats"
"MLJModels"
"BetaML"
"MLJLinearModels"
"LIBSVM"
"EvoTrees"
"NaiveBayes"
"MLJFlux"
"Clustering"
"ParallelKMeans"
"NearestNeighborModels"
"PartialLeastSquaresRegressor"
"LightGBM"
"GLM"
"TSVD"
"MLJText"
"XGBoost"
```

Some multi-paradigm machine learning toolboxes in Julia

- **ScikitLearn.jl** - thin wrapper for Python's scikit-learn
- **AutoMLPipeline.jl** - IBM
- **MLJ.jl**
- FastAI.jl - specific to neural networks (interface to **Flux.jl**)

What?

What's a machine learning toolbox?

What?

What's a machine learning toolbox?

- A toolbox provides a **uniform interface** for **fitting**, **evaluating** and **tuning** machine learning models.
- Provides common **preprocessing tasks** (such as data cleaning and type coercion)
- Allows for model **composition** (e.g., pipelining)

Why?

Why learn the MLJ toolbox?

- Written in Julia (but does wrap non-native models):

Why?

Why learn the MLJ toolbox?

- Written in Julia (but does wrap non-native models):
 - Easy access to core algorithms
 - Easier to customize
 - Easier to add new performant models
 - Greater transparency and reproducibility
 - Better composability with other libraries

Why?

Why learn the MLJ toolbox?

- Written in Julia (but does wrap non-native models):
 - Easy access to core algorithms
 - Easier to customize
 - Easier to add new performant models
 - Greater transparency and reproducibility
 - Better composability with other libraries
- Start-of-the-art **model composition** that plays well with everything else.
- Meta-algorithms (e.g., tuning) are model wrappers

How to entertain your kids at home - the case for MLJ



3 models dinosaurs + 15 user contributed
models

We **support** you in building your snake powered train!



Workshop Overview

- Recap of supervised and unsupervised learning
- Preview of model composition (stacking)
- Part 1 - **Data Representation** + exercises
- Part 2 - **Selecting, Training and Evaluating Models** + exercises
- Break
- Part 3 - **Transformers and Pipelines** + exercises
- Part 4 - **Tuning hyper-parameters** + exercises
- Part 5 - **Advanced model composition** (time permitting)

Supervised Learning

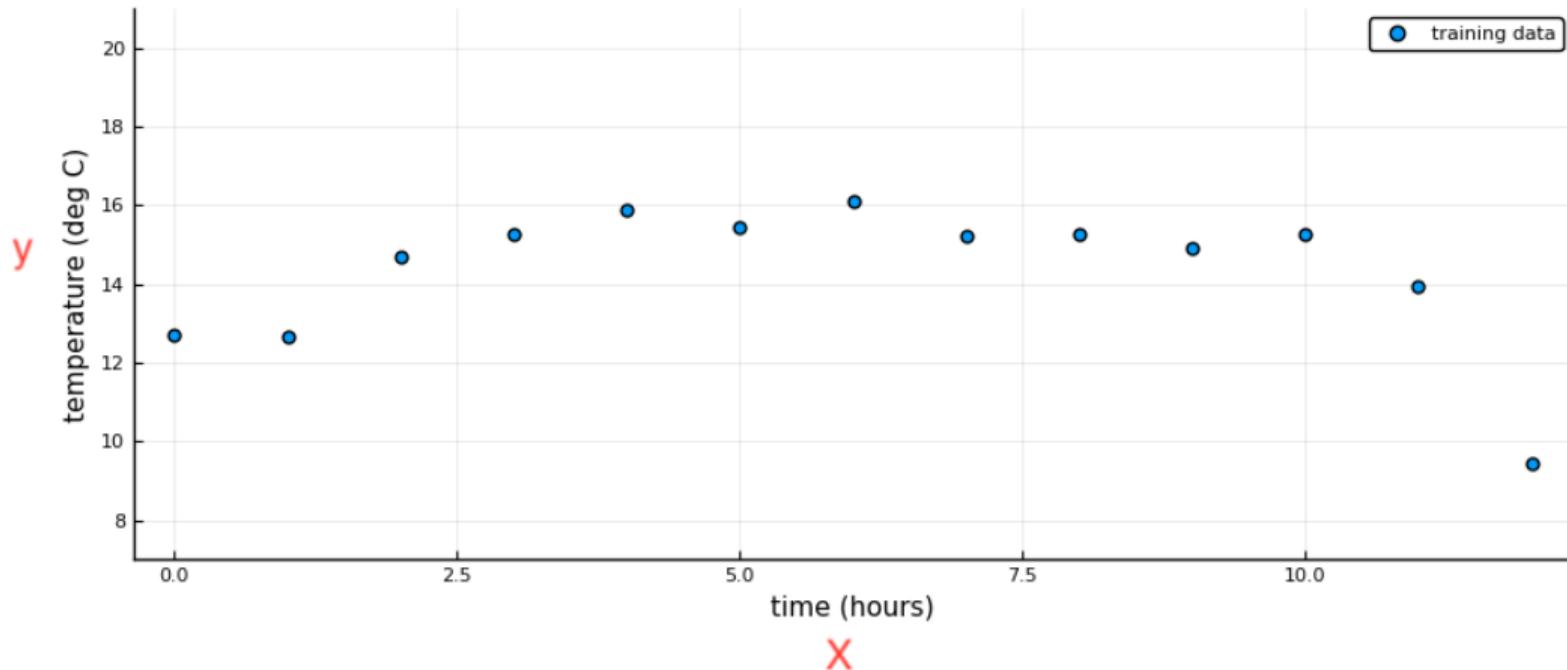
Learning to **predict** some target variable **y** from a knowledge of some other variables **X** (the *input features*).

Supervised Learning

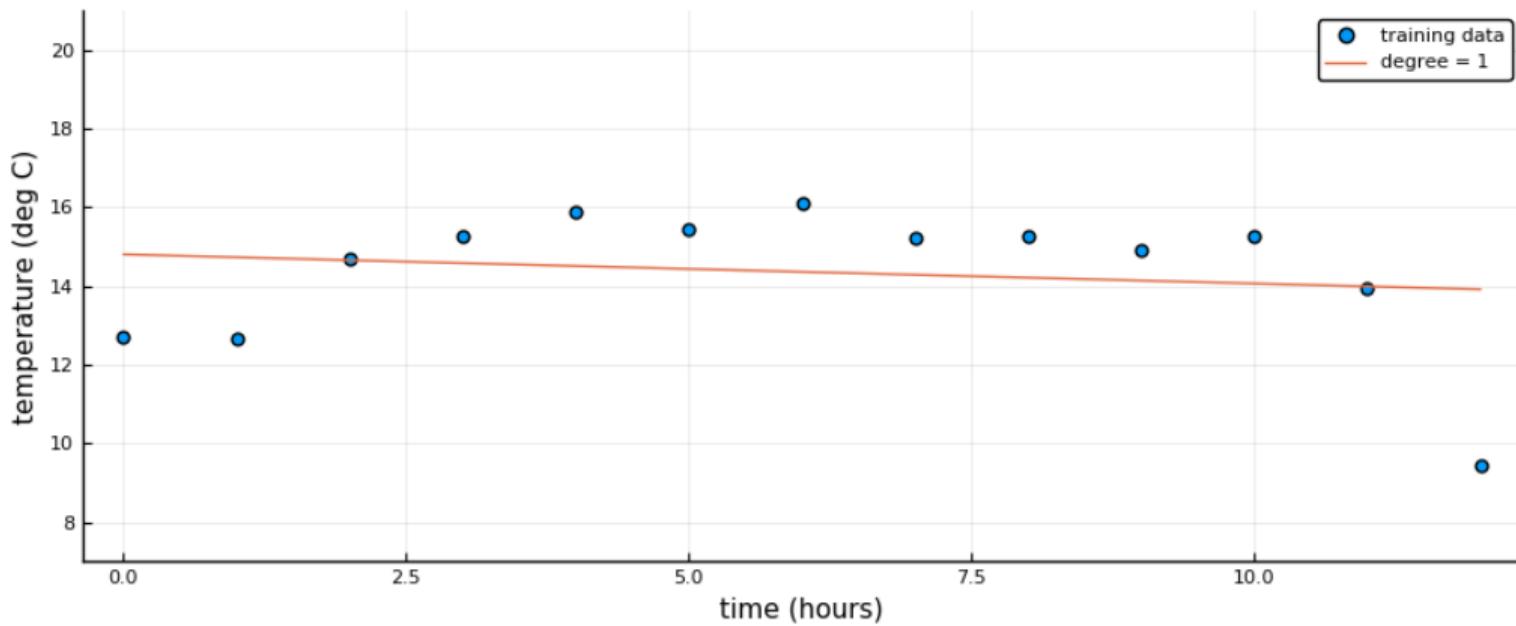
Learning to **predict** some target variable **y** from a knowledge of some other variables **X** (the *input features*).

X		y
time	room	temperature
5	kitchen	18.5
5	bathroom	18.3
5	bedroom_1	18.3
5	living_room	17.4
6	kitchen	16.6
6	bathroom	20.7
6	bedroom_1	18.9
6	living_room	20.2
7	kitchen	20.4
7	bathroom	19.9
7	bedroom_1	16.2
7	living_room	17.3

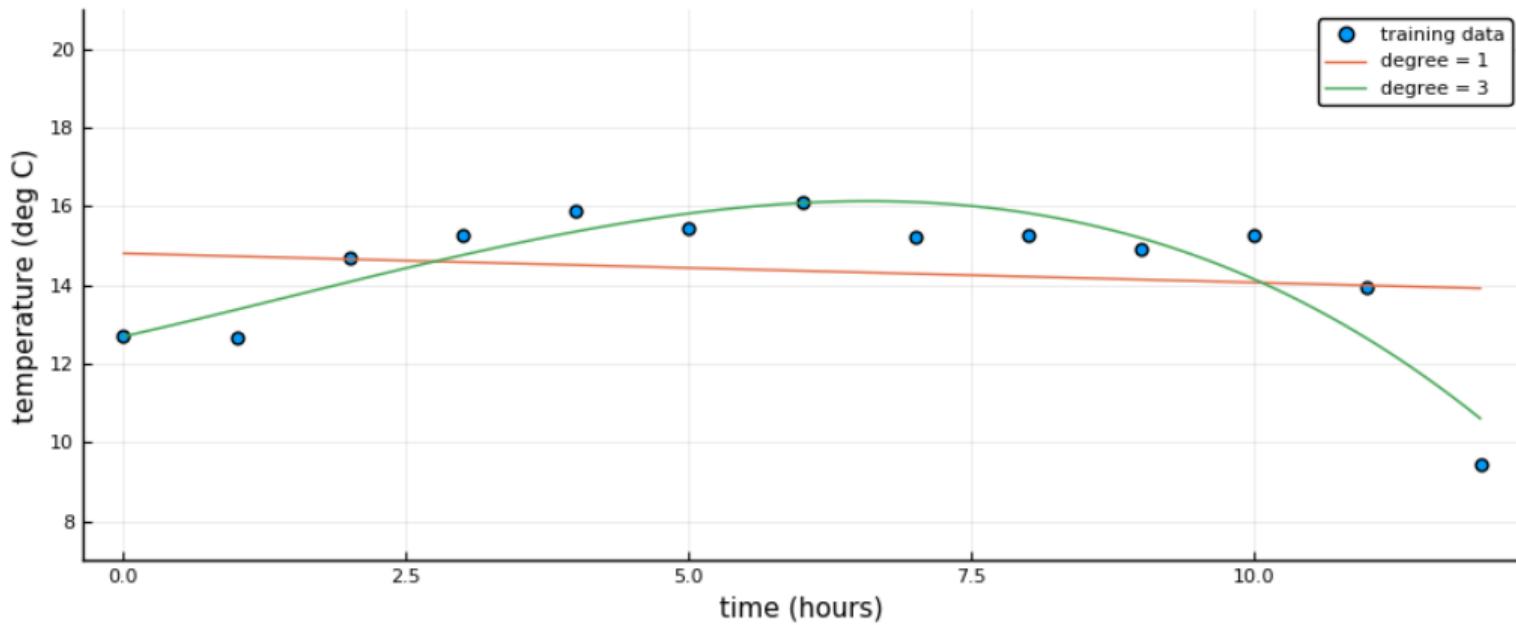
Supervised Learning



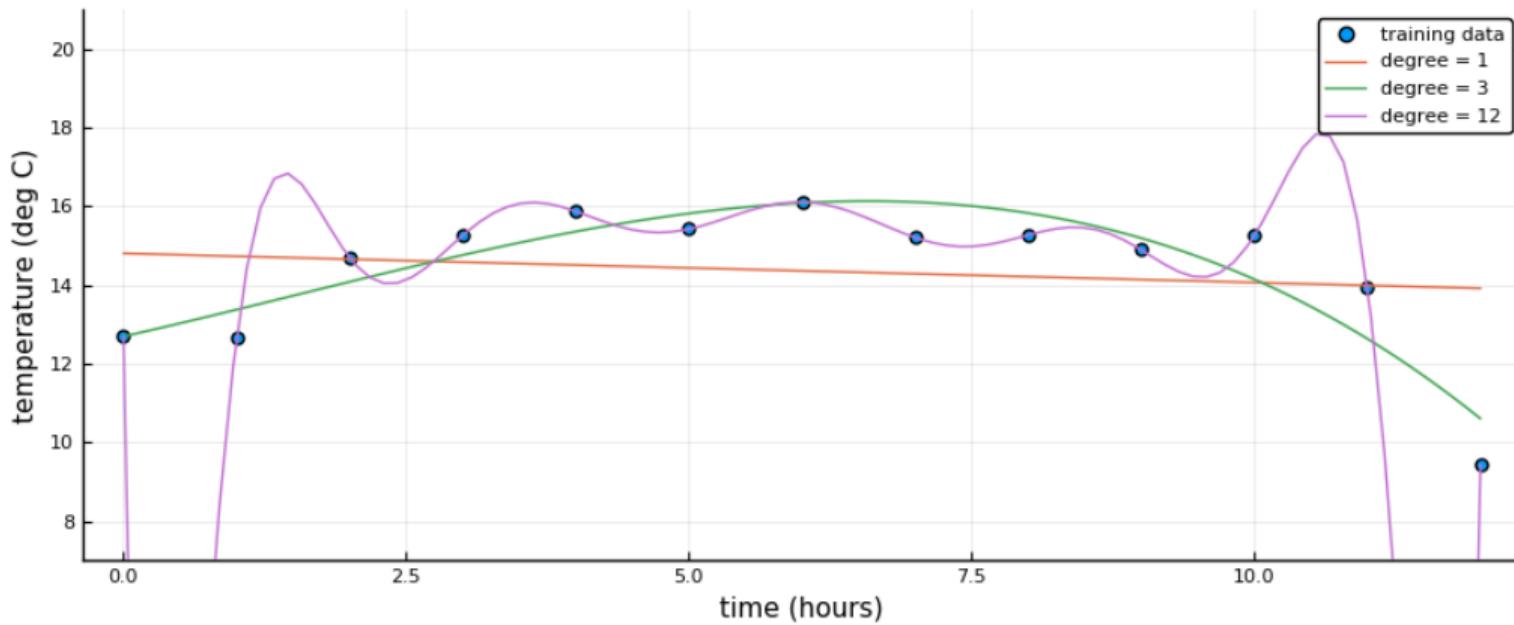
Supervised Learning



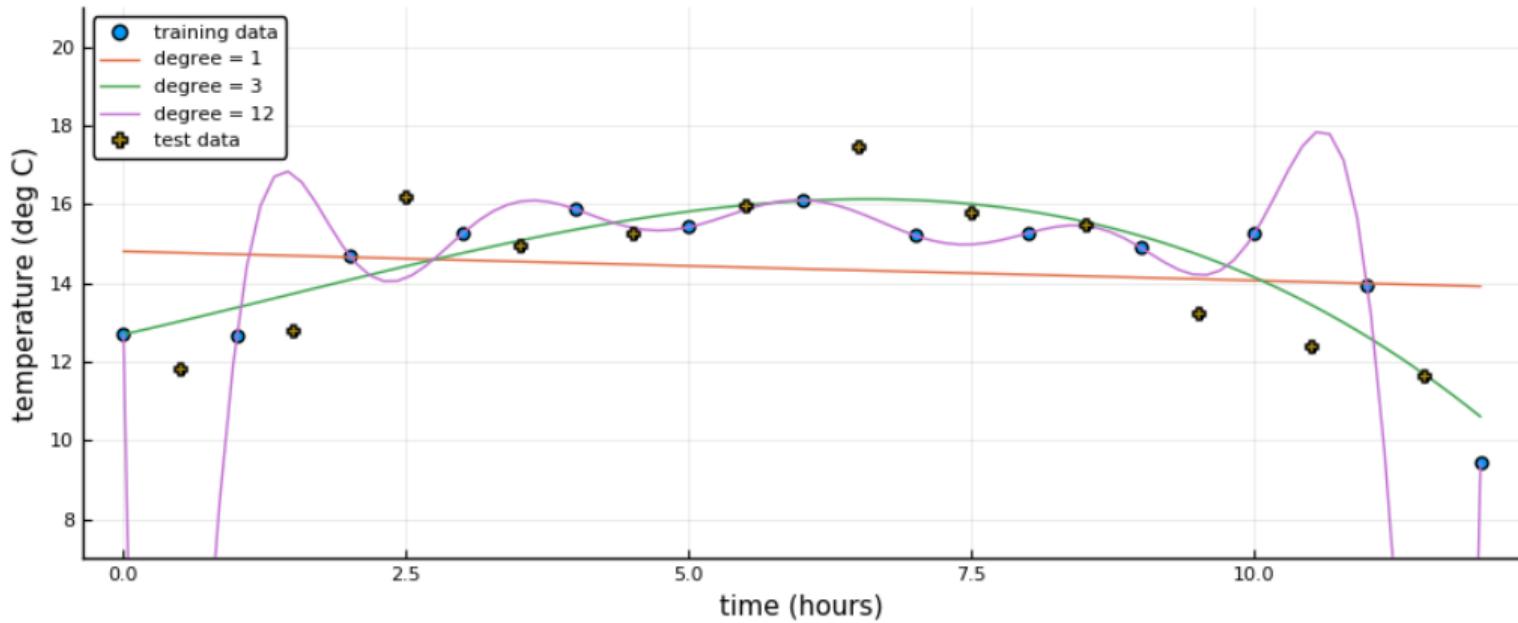
Supervised Learning



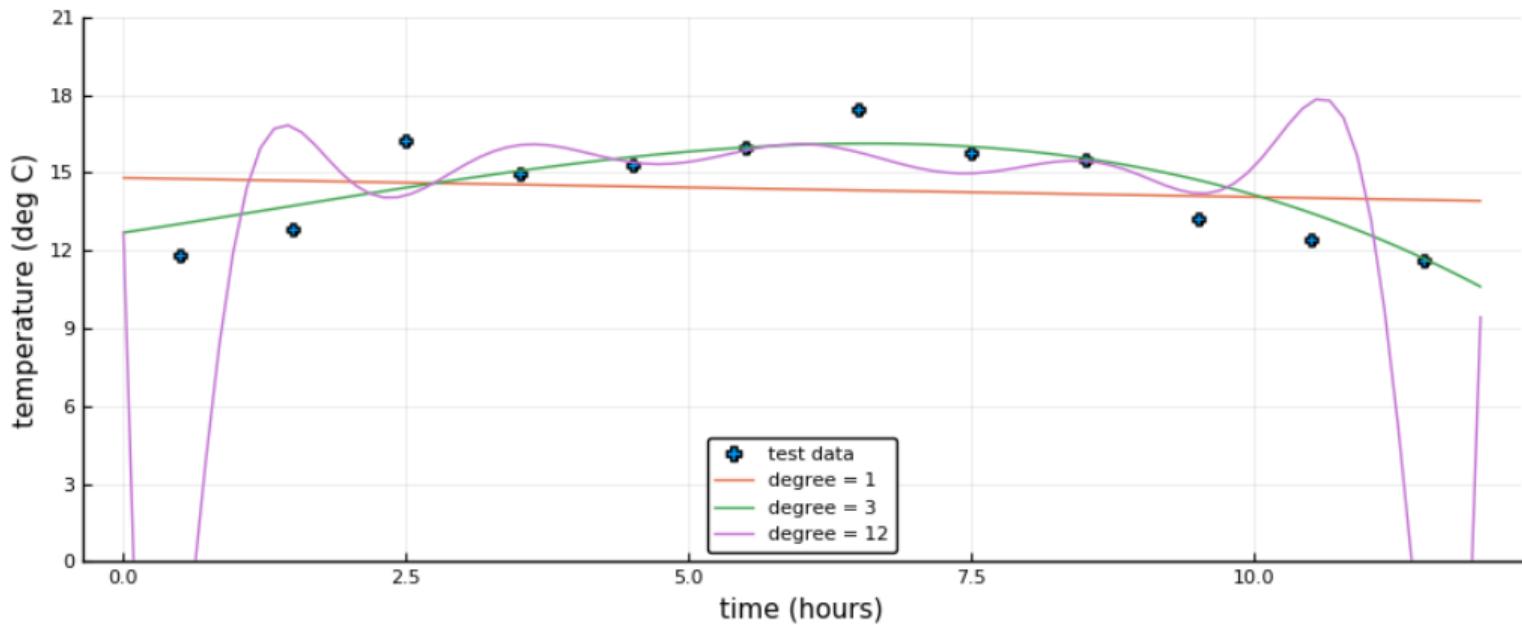
Supervised Learning



Supervised Learning

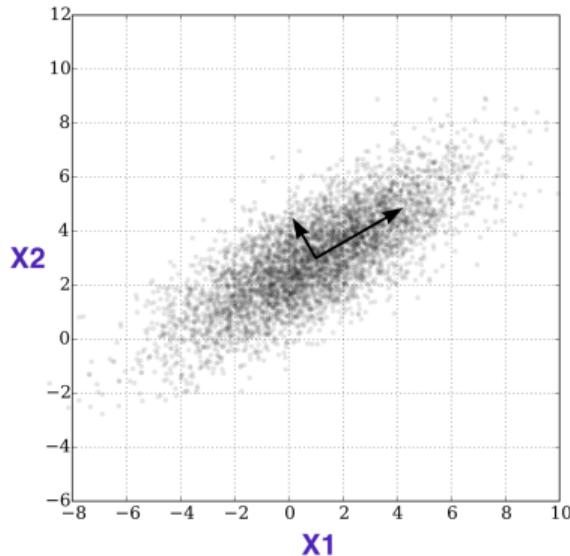


Supervised Learning



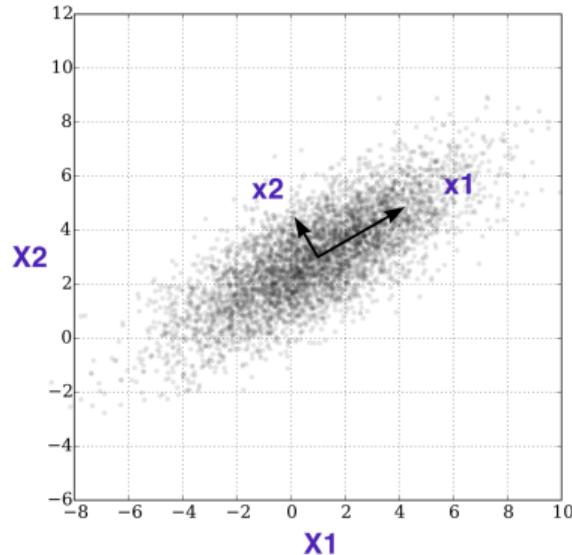
Unsupervised Learning

Learning data **transformations**, e.g., dimension reduction



Unsupervised Learning

Learning data **transformations**, e.g., dimension reduction



Data science competitions (kaggle)

Installing the tutorials

github.com/ablaom/MLJTutorial

turing.ac.uk
@turinginst